



UNIVERSITÀ  
DI TRENTO

# Machine Unlearning

---

Master's Degree in Information Engineering  
Course of Multimedia Data Security

---

Nicola Cappellaro  
Riccardo Zannoni

---

# TABLE OF CONTENTS

---



UNIVERSITÀ  
DI TRENTO

- 1 Introduction
- 2 Project structure
- 3 Implementation
  - 3a Training
  - 3b Poisoning
  - 3c Unlearning
    - c1 Loss function
    - c2 SVD + CGS
    - c3 Gradient update
- 4 Results
- 5 Conclusion

# INTRODUCTION


---



UNIVERSITÀ  
DI TRENTO

## WHAT?

Machine unlearning refers to techniques designed to **selectively remove** the influence of specific data from an already-trained model. The goal is to make the model “forget” specific information **without requiring a full and expensive retraining** from scratch.



## WHY?

Privacy ■

GDPR ■

Data poisoning ■

Right to be forgotten ■

# NN MODEL


---



UNIVERSITÀ  
DI TRENTO

The model is a modified ResNet50 where:

- The first convolutional layer is modified to prevent early spatial downsampling.
- all pretrained weights are kept
- The fully connected layer is replaced to produce a single output



# DATASET

---



UNIVERSITÀ  
DI TRENTO

**Path:** /media/NAS/TrueFake

Only **PreSocial** folder:

- **Real**

FFHQ: 70000 images


FORLAB: 30719 images

- **Fake**

SDXL: 40000 images

GAN2: 40000 images

GAN3: 40000 images



# PROJECT STRUCTURE

- **dataset.py**: dataset loading, preprocessing, and controlled data poisoning.
- **networks.py**: model architecture (modified ResNet50) with optional layer freezing.
- **train.py** : standard training pipeline for the ResNet-based classifier.
- **unlearn\_trueface.py**: implementation of the projected-gradient unlearning procedure.
- **unlearn\_utils.py**: SVD computation, subspace projection utilities, and evaluation helpers.




UNIVERSITÀ  
DI TRENTO

```
├── runs
│   ├── 0_poison
│   ├── 20_poison
│   └── 50_poison
├── dataset.py
├── launcher.py
├── networks.py
├── parser.py
└── train.py
├── unlearn_trueface.py
├── unlearn_utils.py
└── run_unlearning.sh
├── environment.yml
└── README.md
```

# IMPLEMENTATION OVERVIEW:



UNIVERSITÀ  
DI TRENTO



# TRAINING

---



UNIVERSITÀ  
DI TRENTO

## TRAINING PARAMETERS:

- Model nodown
- Freeze ResNet50 weight
- Learning rate 0.0001
- lr\_decay\_epochs 3
- num\_epochs 10
- batch\_size 16
- optimizer: Adam
- loss: BCEWithLogitsLoss


## DATA AUGMENTATION

- Image normalization,
  - resize
  - crops
  - blur
  - JPEG compression
- improve robustness and mimic real world distortions.

# TRAINING TRAINING LOOP



UNIVERSITÀ  
DI TRENTO





We apply poison by flipping the dataset labels with a given  
**poison rate** (e.g., 10%, 20%)

**Selection:** `random.sample()` on dataset indices

Labels are **flipped** as follows:  
Real (0) → Fake (1)  
Fake (1) → Real (0)

**Tracking:** poisoned indices saved to  
runs/poison\_info/poison\_xx.pkl for unlearning phase

# UNLEARNING

---



UNIVERSITÀ  
DI TRENTO

## DETAILED OVERVIEW:

Based on **Projected Gradient Unlearning** (PGU)

### Goal:

- Reduce the model's confidence on the data that must be "forgotten"
- Preserve performance on the remaining clean data

### Three **key components**:

- Unlearning loss
- SVD and CGS Construction
- Projected gradient update

# LOSS FUNCTION

---



UNIVERSITÀ  
DI TRENTO

**Loss 1:** reduce the model's confidence on the poisoned label:

$$loss1 = -\log(1 - |p - (1 - label)| - offset)$$

**Loss 2:** increase the entropy of the predictions:

$$loss2 = -H(p), \text{ where}$$

$$H(p) = -p \log(p) - (1 - p) \log(1 - p)$$

**Total loss:** weighted sum of the two terms:

$$L = w1 \cdot loss1 + w2 \cdot loss2$$

# SVD + CGS



UNIVERSITÀ  
DI TRENTO

We want to **remove** the influence of poisoned data without damaging the model



To understand **which directions** represent useful knowledge, we compute the **SVD** of clean-data activations.



We keep only the **most important directions** (e.g., 95% variance): this forms the clean-data subspace  $U_k$



This matrix is used to **project gradients** onto the clean-data subspace.

$$g_{proj} = Pg$$



We then build the projection matrix using **CGS**:

$$P = U_k U_k^T$$

# GRADIENT UPDATE



UNIVERSITÀ  
DI TRENTO

After computing the unlearning loss and projecting the gradient, we update **only the safe part** of the model. The projected gradient  $g_{\text{proj}}$  is applied to the parameters:

$$\theta \leftarrow \theta - \eta g_{\text{proj}}$$

This ensures the update:

1. **removes** the influence of poisoned data
2. **preserves** the clean-data subspace learned by the model, avoiding catastrophic forgetting.

Only the final fully connected layer is updated when freeze=True.

→ The model is updated using only the gradient components that do not interfere with clean-data knowledge.

# RESULTS

---



UNIVERSITÀ  
DI TRENTO

## Performance 0% poison

Accuracy: 97.79%  
Precision: 98.03%  
Recall: 97.91%  
F1 Score: 97.97%

## Performance 20% poison

Accuracy: 96.39%  
Precision: 94.17%  
Recall: 99.54%  
F1 Score: 96.78%

## Performance 50% poison

Accuracy: 42.62%  
Precision: 46.19%  
Recall: 31.74%  
F1 Score: 37.62%

## Performance 20% unlearn

Initial Clean Train Acc: 96.48%  
Final Clean Train Acc: 97.52%  
Clean Acc: +1.04%  
  
Initial Poison Train Acc: 99.48%  
Final Poison Train Acc: 98.73%  
Poison Acc: -0.75%

Initial Test Acc: 96.67%  
Final Test Acc: 97.64%  
Test Acc Change: +0.97%

## Performance 50% unlearn

Initial Clean Train Acc: 42.20%  
Final Clean Train Acc: 48.74%  
Clean Acc: +6.54%  
  
Initial Poison Train Acc: 55.04%  
Final Poison Train Acc: 0.00%  
Poison Acc: -55.04%

Initial Test Acc: 42.62%  
Final Test Acc: 45.47%  
Test Acc Change: +2.86%

# CONCLUSION

---

## UNLEARNING 20% POISON

1. The model with 20% poison was already excellent, almost identical to the clean model.
2. The amount of harmful information is small → unlearning only slightly changes the behavior.

Unlearning removes the “poisoned” component, but since the model was already very good, this results in only a slight improvement (~1%).

## UNLEARNING 50% POISON

1. The model is severely corrupted: accuracy collapses to ~42%, so unlearning has much more to remove.
2. The unlearning process successfully eliminates this poisoned influence (poison accuracy drops to 0%), and the model partially recovers: clean accuracy improves (+6.5%)

Although performance cannot return to clean-model levels, this is expected: the goal of PGU is to forget the poisoned data without harming the clean data.



UNIVERSITÀ  
DI TRENTO



UNIVERSITÀ  
DI TRENTO

**THANK YOU FOR  
YOUR ATTENTION**

---

Nicola Cappellaro – [nicola.cappellaro@studenti.unitn.it](mailto:nicola.cappellaro@studenti.unitn.it)  
Riccardo Zannoni – [riccardo.zannoni@studenti.unitn.it](mailto:riccardo.zannoni@studenti.unitn.it)

---